Change 1: Decoupling hidden size and head number

Adjustable	Fixed ❌	→	Adjustable	Adjustable ✅
KV Cache	Attention FLOPs		KV Cache	Attention FLOPs
Ordinary GQA				*Flexible adjustment of time-variant compute*

Change 2: Joint optimization of GQA configuration and model size

Adjustable	Fixed ❌	→	Adjustable	Adjustable ✅
Time-variant memory/FLOPs	Time-invariant memory/FLOPs		Time-variant memory/FLOPs	Time-invariant memory/FLOPs
*Faster but **worse***			*Faster and **better***	

Result Inference costs (@128K)

