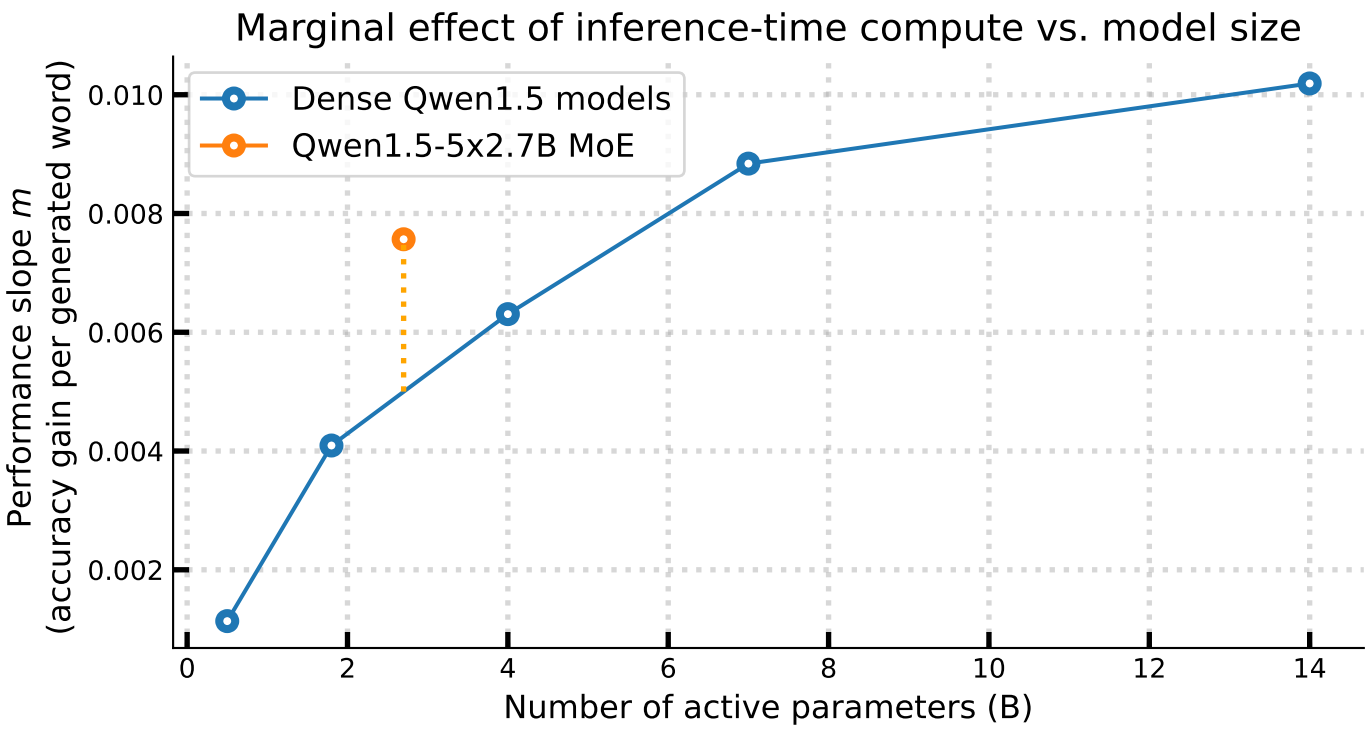
Marginal effect of inference-time compute vs. model size
Dense Qwen1.5 models
Qwen1.5-5x2.7B MoE
Performance slope m
(accuracy gain per generated word)
0.010
0.008
0.006
0.004
0.002
0
2
4
6
8
10
12
14
Number of active parameters (B)