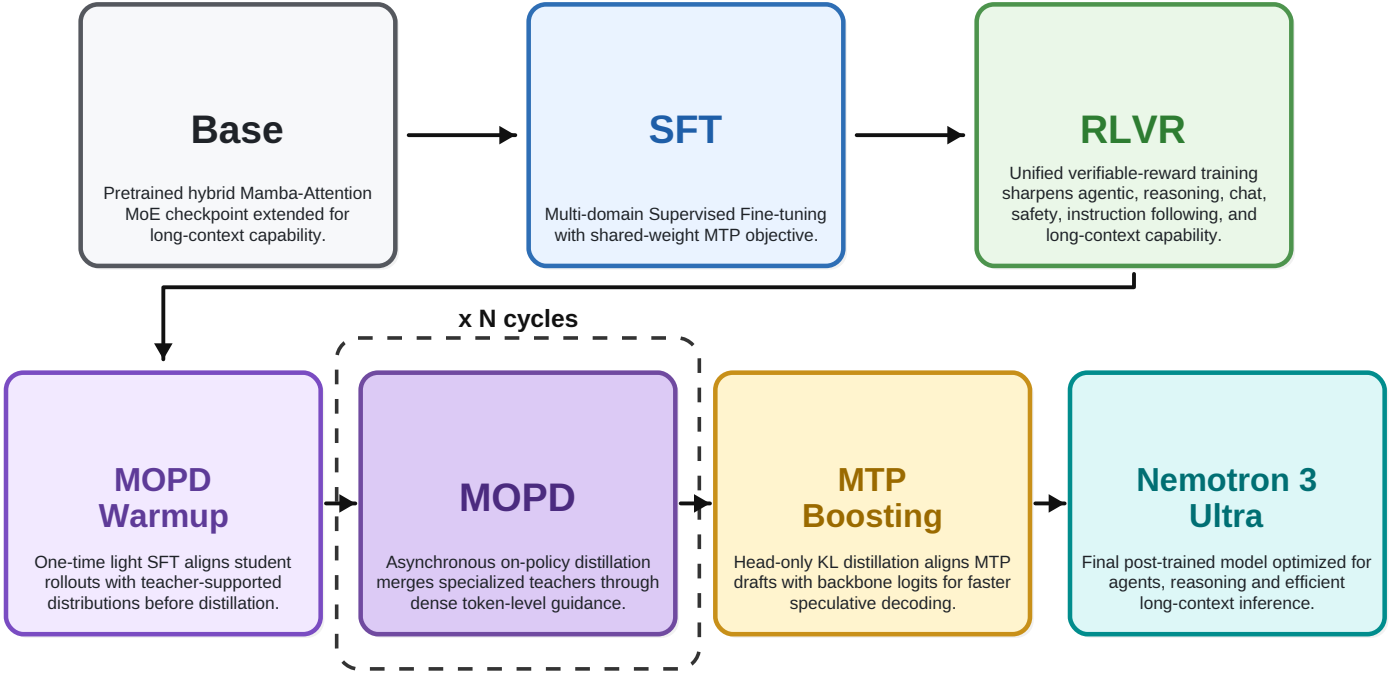
Base
Pretrained hybrid Mamba-Attention MoE checkpoint extended for long-context capability.
SFT
Multi-domain Supervised Fine-tuning with shared-weight MTP objective.
RLVR
Unified verifiable-reward training sharpens agentic, reasoning, chat, safety, instruction following, and long-context capability.
x N cycles
MOPD Warmup
One-time light SFT aligns student rollouts with teacher-supported distributions before distillation.
MOPD
Asynchronous on-policy distillation merges specialized teachers through dense token-level guidance.
MTP Boosting
Head-only KL distillation aligns MTP drafts with backbone logits for faster speculative decoding.
Nemotron 3 Ultra
Final post-trained model optimized for agents, reasoning and efficient long-context inference.